

Rook

前言

Rook是一个自我管理的分布式存储编排系统，它本身并不是存储系统，在存储和 k8s之前搭建了一个桥梁，使存储系统的搭建或者维护变得特别简单，Rook将分布式存储系统转变为自我管理、自我扩展、自我修复的存储服务。它让一些存储的操作，比如部署、配置、扩容、升级、迁移、灾难恢复、监视和资源管理变得自动化，无需人工处理。并且 Rook支持CSI，可以利用CSI做一些PVC的快照、扩容、克隆等操作。

我们不生产水，我们只是大自然的搬运工

Rook

前言

Rook 架构

组件介绍：

基础环境

Rook 部署

- 1) 克隆 Rook 仓库
- 2) 下载镜像
- 3) 修改配置文件

ceph 部署集群

- 1) 创建Ceph集群
- 2) 安装 ceph snapshot 控制器
- 3) 安装 ceph 客户端工具
- 4) 安装Ceph dashboard

ceph 块存储的使用

- 1) 创建 StorageClass 和 ceph 的存储池
- 2) 挂载测试
- 3) StatefulSet 动态存储

共享文件系统的使用

- 1) 创建共享类型的文件系统和 StorageClass
- 2) nginx 挂载测试
- 3) 验证效果

PVC 扩容

- 1) 扩容文件共享型 PVC
- 2) 扩容块存储

PVC 快照

- 1) 创建 snapshotClass
- 2) 创建快照
- 3) 通过快照创建 PVC
- 4) 数据校验
- 5) 文件共享类型快照

PVC 克隆


测试数据清理

- 1.清理挂载了 PVC 的 Pod 和 Deployment
- 2.清理 PVC
- 3.清理快照
- 4.清理存储池
- 5.清理 StorageClass
- 6.清理 Ceph 集群

- 7.删除 Rook 资源
- 8.处理卡住资源（如有）
- 9.清理数据目录和磁盘
- 10.清理OSD 所使用磁盘

Rook 架构

Rook 架构



组件介绍:

核心Pod功能说明

- `rook-discover`: Rook设备发现守护进程，**自动扫描节点上的存储设备**（如磁盘、SSD），并将信息上报给Operator。
- `rook-ceph-mon`: Ceph Monitor服务，**维护集群元数据**（如OSD映射、PG状态），**确保集群一致性**。
- `rook-ceph-mgr`: Ceph Manager服务，**提供管理接口和监控指标**（Dashboard、Prometheus指标）。
- `rook-ceph-osd`: Ceph OSD服务，**实际存储数据的守护进程**，每个OSD对应一个物理设备。
- `rook-ceph-crashcollector`: 崩溃日志收集器，**自动收集OSD/Mon故障时的诊断信息**。
- `rook-ceph-exporter`: 指标导出器，**将Ceph集群性能数据导出给Prometheus**。

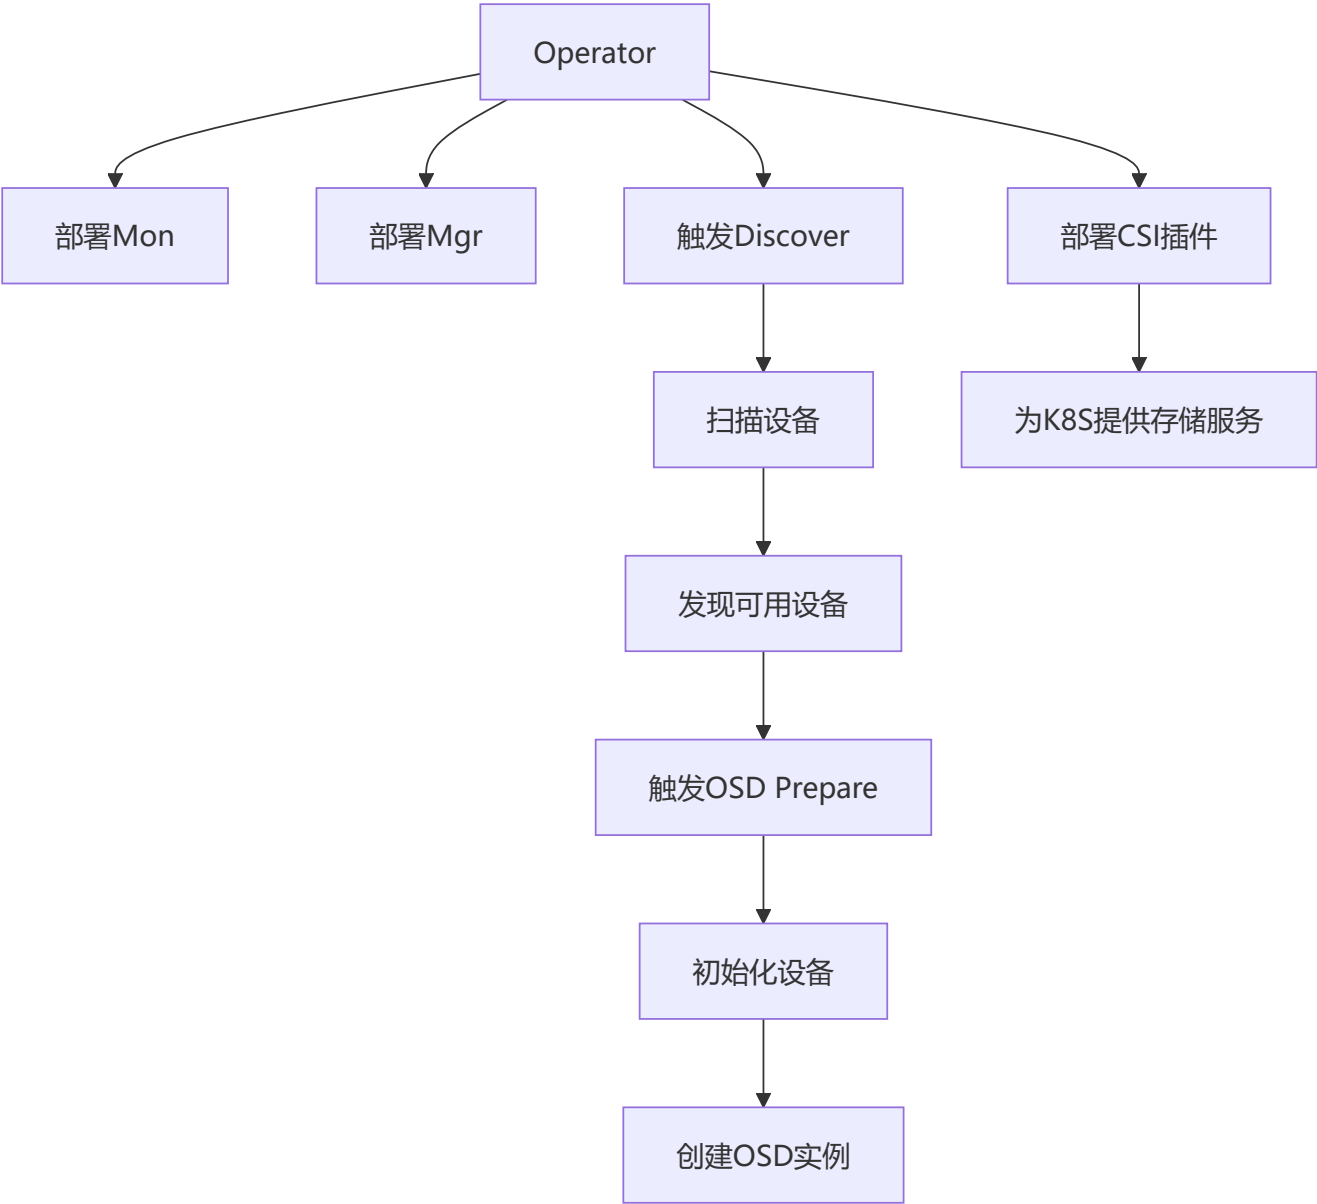
CSI相关组件说明

- `csi-rbdplugin`: RBD块存储插件，**支持动态创建/挂载Ceph RBD卷**（适用于数据库等场景）。
- `csi-cephfsplugin`: CephFS文件存储插件，**支持动态创建/挂载CephFS卷**（适用于共享文件存储）。
- `csi-*provisioner`: 存储供应控制器，**处理存储卷的生命周期管理**（创建/扩容/删除）。

运维工具类说明

- `rook-ceph-operator`: 集群管理核心，**协调所有资源部署和配置更新**。

- `rook-ceph-tools`: 工具Pod (需手动创建), 提供 `ceph` 命令行工具, 用于集群调试。



基础环境

name	IP	磁盘	ceph
k8s-master01	192.168.0.105	nvme0n1	
k8s-master02	192.168.0.106	nvme0n1	
k8s-master03	192.168.0.107	nvme0n1、nvme0n2	ceph安装所在节点
k8s-work01	192.168.0.115	nvme0n1、nvme0n2	ceph安装所在节点
k8s-work02	192.168.0.116	nvme0n1、nvme0n2	ceph安装所在节点

Rook 部署

集群运行后, 应用程序可以使用块、对象或文件存储。

1) 克隆 Rook 仓库

```
1 | git clone --single-branch --branch v1.16.5 https://github.com/rook/rook.git
```

2) 下载镜像

按照利用阿里云下载国外镜像文档进行下载

```
1 | ctr images pull registry.cn-hangzhou.aliyuncs.com/xusx/images:3.13.0
2 | ctr images tag registry.cn-hangzhou.aliyuncs.com/xusx/images:3.13.0
   | quay.io/cephcsi/cephcsi:v3.13.0
3 |
4 |
5 | ctr images pull registry.cn-hangzhou.aliyuncs.com/xusx/images:2.13.0
6 | ctr images tag registry.cn-hangzhou.aliyuncs.com/xusx/images:2.13.0
   | registry.k8s.io/sig-storage/csi-node-driver-registrar:v2.13.0
7 |
8 |
9 | ctr images pull registry.cn-hangzhou.aliyuncs.com/xusx/images:1.13.1
10 | ctr images tag registry.cn-hangzhou.aliyuncs.com/xusx/images:1.13.1
    | registry.k8s.io/sig-storage/csi-resizer:v1.13.1
11 |
12 |
13 | ctr images pull registry.cn-hangzhou.aliyuncs.com/xusx/images:5.1.0
14 | ctr images tag registry.cn-hangzhou.aliyuncs.com/xusx/images:5.1.0 registry.k8s.io/sig-
    | storage/csi-provisioner:v5.1.0
15 |
16 |
17 | ctr images pull registry.cn-hangzhou.aliyuncs.com/xusx/images:8.2.0
18 | ctr images tag registry.cn-hangzhou.aliyuncs.com/xusx/images:8.2.0 registry.k8s.io/sig-
    | storage/csi-snapshotter:v8.2.0
19 |
20 |
21 | ctr images pull registry.cn-hangzhou.aliyuncs.com/xusx/images:4.8.0
22 | ctr images tag registry.cn-hangzhou.aliyuncs.com/xusx/images:4.8.0 registry.k8s.io/sig-
    | storage/csi-attacher:v4.8.0
23 |
24 |
25 | ctr images pull registry.cn-hangzhou.aliyuncs.com/xusx/images:1.16.5
26 | ctr images tag registry.cn-hangzhou.aliyuncs.com/xusx/images:1.16.5
    | docker.io/rook/ceph:v1.16.5
27 |
28 | ctr images pull registry.cn-hangzhou.aliyuncs.com/xusx/images:2.13.0
29 | ctr images tag registry.cn-hangzhou.aliyuncs.com/xusx/images:2.13.0
    | registry.k8s.io/sig-storage/csi-node-driver-registrar:v2.13.0
```

3) 修改配置文件

修改 Rook CSI 镜像地址，原本的地址可能是 k8s.io 的镜像，但是无法被国内访问，所以需要同步gcr的镜像到阿里云镜像仓库

operator 文件，新版本 rook 默认关闭了自动发现容器的部署，可以找到 ROOK_ENABLE_DISCOVERY_DAEMON 改成 true

```
1 sed -i -E 's/(ROOK_ENABLE_DISCOVERY_DAEMON:\s*)"false"/\1"true"/g' operator.yaml
2 # -E: 启用扩展正则表达式（支持 \s* 匹配任意数量空格）
3 # \s*: 匹配键值之间的任意数量空格。
4 # \1: 保留原键名和冒号后的格式，仅替换值部分
5
6 cd rook/deploy/examples
7 kubectl create namespace rook-ceph
8 kubectl create -f crds.yaml -f common.yaml -f operator.yaml
9
10 [root@K8S-Master01 examples]# kubectl -n rook-ceph get pod
11 NAME                                READY   STATUS    RESTARTS   AGE
12 rook-ceph-operator-67944bdfcc-b7r79 1/1     Running   0           35s
13 rook-discover-l4cw8                 1/1     Running   0           33s
14 rook-discover-qf5z8                 1/1     Running   0           33s
15 [root@K8S-Master01 examples]# kubectl get crd |grep rook
16 cephblockpoolradosnamespaces.ceph.rook.io      2025-03-17T09:03:29Z
17 cephblockpools.ceph.rook.io                    2025-03-17T09:03:29Z
18 cephbucketnotifications.ceph.rook.io           2025-03-17T09:03:29Z
19 cephbuckettopics.ceph.rook.io                  2025-03-17T09:03:29Z
20 cephclients.ceph.rook.io                       2025-03-17T09:03:29Z
21 cephclusters.ceph.rook.io                      2025-03-17T09:03:29Z
22 cephcosidrivers.ceph.rook.io                   2025-03-17T09:03:29Z
23 cephfilesystemmirrors.ceph.rook.io             2025-03-17T09:03:29Z
24 cephfilesystems.ceph.rook.io                   2025-03-17T09:03:29Z
25 cephfilesystemsubvolumegroups.ceph.rook.io     2025-03-17T09:03:29Z
26 cephnfses.ceph.rook.io                        2025-03-17T09:03:30Z
27 cephobjectrealms.ceph.rook.io                  2025-03-17T09:03:30Z
28 cephobjectstores.ceph.rook.io                  2025-03-17T09:03:30Z
29 cephobjectstoreusers.ceph.rook.io              2025-03-17T09:03:30Z
30 cephobjectzonegroups.ceph.rook.io              2025-03-17T09:03:30Z
31 cephobjectzones.ceph.rook.io                   2025-03-17T09:03:30Z
32 cephrbdmirrors.ceph.rook.io                   2025-03-17T09:03:30Z
```

ceph 部署集群

注意：新版必须采用裸盘，即未格式化的磁盘。其中 k8s-master03 k8s-node01 node02 有新加的一个磁盘，可以通过 `lsblk -f` 查看新添加的磁盘名称。建议最少三个节点，否则后面的试验可能会出现问

```

1 [root@K8S-Master03 ~]# lsblk
2 NAME          MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
3 sr0            11:0    1 1024M  0 rom
4 nvme0n1        259:0    0  100G  0 disk
5 └─nvme0n1p1    259:1    0    1G  0 part /boot
6 └─nvme0n1p2    259:2    0   99G  0 part
7   └─r1-root    253:0    0   99G  0 lvm  /
8 nvme0n2        259:3    0   10G  0 disk

```

#新加磁盘，用作ceph OSD

1) 创建Ceph集群

```

storage: # cluster level storage configuration and selection
  useAllNodes: false # 在所有节点上进行部署
  useAllDevices: false # 使用所有的磁盘设备部署
  #deviceFilter:
  config:
    # crushRoot: "custom-root" # specify a non-default root label for the CRUSH map
    # metadataDevice: "md0" # specify a non-rotational storage so ceph-volume will use it as block db device of bluestore.
    # databaseSizeMB: "1024" # uncomment if the disks are smaller than 100 GB
    # osdsPerDevice: "1" # this value can be overridden at the node or device level
    # encryptedDevice: "true" # the default value for this option is "false"
    # deviceClass: "myclass" # specify a device class for OSDs in the cluster
    allowDeviceClassUpdate: false # whether to allow changing the device class of an OSD after it is created
    allowOsdCrushWeightUpdate: false # whether to allow resizing the OSD crush weight after osd pvc is increased
    # Individual nodes and their config can be specified as well, but 'useAllNodes' above must be set to false. Then, only the named
    # nodes below will be used as storage resources. Each node's 'name' field should match their 'kubernetes.io/hostname' label.
    # nodes:
    #   - name: "172.17.4.201"
    #   devices: # specific devices to use for storage can be specified for each node
    #   #   - name: "sdb"
    #   #   - name: "nvme01" # multiple osds can be created on high performance devices
    #   config:
    #     osdsPerDevice: "5"
    #   - name: "/dev/disk/by-id/ata-ST4000DM004-XXXX" # devices can be specified using full udev paths
    #   config: # configuration can be specified at the node level which overrides the cluster level config
    #   - name: "172.17.4.301"
    #   deviceFilter: "^sd."
    # Whether to always schedule OSD pods on nodes declared explicitly in the "nodes" section, even if they are
    # temporarily not schedulable. If set to true, consider adding placement tolerations for unschedulable nodes.
    nodes:
    - name: "K8S-Master03"
      devices:
      - name: "nvme0n2"
    - name: "K8S-Work01"
      devices:
      - name: "nvme0n2"
    - name: "K8S-Work02"
      devices:
      - name: "nvme0n2"

```

输入要部署的节点，和磁盘设备名称

```

# To understand Rook's upgrade process of ceph, read https://rook.io/docs/rook/latest/ceph-upgrade.html#ceph-version-upgrades
skipUpgradeChecks: false # 是否跳过升级检查
# Whether or not continue if PGs are not clean during an upgrade
continueUpgradeAfterChecksEvenIfNotHealthy: false
# WaitTimeForHealthyOSDMinutes defines the time (in minutes) the operator would wait before an OSD can be stopped for upgrade or restart.
# If the timeout exceeds and OSD is not ok to stop, then the operator would skip upgrade for the current OSD and proceed with the next one
# If 'continueUpgradeAfterChecksEvenIfNotHealthy' is 'false', If 'continueUpgradeAfterChecksEvenIfNotHealthy' is 'true', then operator would
# continue with the upgrade of an OSD even if its not ok to stop after the timeout. This timeout won't be applied if 'skipUpgradeChecks' is 'true'.
# The default wait timeout is 10 minutes.
waitTimeoutForHealthyOSDInMinutes: 10
# Whether or not requires PGs are clean before an OSD upgrade. If set to 'true' OSD upgrade process won't start until PGs are healthy.
# This configuration will be ignored if 'skipUpgradeChecks' is 'true'.
# Default is false.
upgradeOSDRequiresHealthyPGs: false
mon:
  # Set the number of mons to be started. Generally recommended to be 3.
  # For highest availability, an odd number of mons should be specified.
  count: 3 # Monitoring副本数量, 生产>=3
  # The mons should be on unique nodes. For production, at least 3 nodes are recommended for this reason.
  # Mons should only be allowed on the same node for test environments where data loss is acceptable.
  allowMultiplePerNode: false # 是否允许副本都部署在同一个节点上
mgr:
  # When higher availability of the mgr is needed, increase the count to 2.
  # In that case, one mgr will be active and one in standby. When Ceph updates which
  # mgr is active, Rook will update the mgr services to match the active mgr.
  count: 2 # MGR副本数量, 生产>=2
  allowMultiplePerNode: false # 是否允许副本都部署在同一个节点上
modules:
  # List of modules to optionally enable or disable.
  # Note the "dashboard" and "monitoring" modules are already configured by other settings in the cluster CR.
  - name: rook
    enabled: true
# enable the ceph dashboard for viewing cluster status
dashboard:
  enabled: true # dashboard是否打开
  # serve the dashboard under a subpath (useful when you are accessing the dashboard via a reverse proxy)
  # urlPrefix: /ceph-dashboard
  # serve the dashboard at the given port.
  # port: 8443
  # serve the dashboard using SSL
  ssl: true # 是否启用ssl, 我们是测试环境 要改为 false, 否正可能无法访问
  # The url of the Prometheus instance
  # prometheusEndpoint: <protocol>://<prometheus-host>:<port>
  # Whether SSL should be verified if the Prometheus server is using https
  # prometheusEndpointSSLVerify: false
# enable prometheus alerting for cluster
monitoring:
  # requires Prometheus to be pre-installed
  enabled: false # 是否开启Prometheus监控

```

```

allowUninstallWithVolumes: false
# To control where various services will be scheduled by kubernetes, use the placement configuration sections below.
# The example under 'all' would have all services scheduled on kubernetes nodes labeled with 'role=storage-node' and
# tolerate taints with a key of 'storage-node'.
# placement:
#   all:
#     nodeAffinity:
#       requiredDuringSchedulingIgnoredDuringExecution:
#         nodeSelectorTerms:
#           - matchExpressions:
#             - key: role
#               operator: In
#               values:
#                 - storage-node
#     podAffinity:
#     podAntiAffinity:
#     topologySpreadConstraints:
#     tolerations:
#       - key: storage-node
#         operator: Exists
# The above placement information can also be specified for mon, osd, and mgr components
# mon:
# Monitor deployments may contain an anti-affinity rule for avoiding monitor
# collocation on the same node. This is a required rule when host network is used
# or when AllowMultiplePerNode is false. Otherwise this anti-affinity rule is a
# preferred rule with weight: 50.
# osd:
#   prepareosd:
#   mgr:
#   cleanup:
# annotations:
#   all:
#   mon:
#   mgr:
#   osd:
#   exporter:
#   crashcollector:
#   cleanup:
#   prepareosd:
# cmdreporter is for jobs to detect ceph and csi versions, and check network status
# cmdreporter:
# clusterMetadata annotations will be applied to only 'rook-ceph-mon-endpoints' configmap and the 'rook-ceph-mon' and 'rook-ceph-admin-keyring' secrets.
# And clusterMetadata annotations will not be merged with 'all' annotations.
# clusterMetadata:
#   kubed.appscode.com/sync: "true"
# If no mgr annotations are set, prometheus scrape annotations will be set by default.
# mgr:
# labels:
#   all:
-- INSERT --

```

如果要部署限制的机器，在这里配置反亲和力和

216,13 53%

```

1 # 修改上述配置
2 vim rook/deploy/examples/cluster.yaml
3
4 # 按照使用阿里云下载外网镜像方法下载: quay.io/ceph/ceph:v19.2.1
5 ctr images pull registry.cn-hangzhou.aliyuncs.com/xusx/images:19.2.1
6 ctr images tag registry.cn-hangzhou.aliyuncs.com/xusx/images:19.2.1
  quay.io/ceph/ceph:v19.2.1
7
8 kubectl create -f cluster.yaml
9 kubectl -n rook-ceph get pod
10 # 查看集群健康
11 kubectl -n rook-ceph get cephcluster

```

```

[root@k8s-master01 examples]# kubectl -n rook-ceph get cephcluster
NAME          DATADIRHOSTPATH  MONCOUNT  AGE    PHASE  MESSAGE                                     HEALTH  EXTERNAL  FSID
rook-ceph    /var/lib/rook    3          9m47s  Ready  Cluster created successfully              HEALTH_OK  EXTERNAL  dd96e25d-c610-4fc2-a0aa-41bf4cd9d750
[root@k8s-master01 examples]# kubectl -n rook-ceph get pods
NAME                                     READY   STATUS    RESTARTS   AGE
csi-cephfsplugin-9zxwx                  3/3     Running   0           9m51s
csi-cephfsplugin-jdndt                  3/3     Running   0           9m51s
csi-cephfsplugin-js62l                  3/3     Running   0           9m51s
csi-cephfsplugin-provisioner-5f589446b4-fdf4h  6/6     Running   0           9m51s
csi-cephfsplugin-provisioner-5f589446b4-vds6n  6/6     Running   0           9m51s
csi-cephfsplugin-pvqxb                  3/3     Running   0           9m51s
csi-rbdplugin-kgdde                     3/3     Running   0           9m52s
csi-rbdplugin-provisioner-bb8f8f69d-6wl8d     6/6     Running   0           9m51s
csi-rbdplugin-provisioner-bb8f8f69d-hl4l7     6/6     Running   0           9m51s
csi-rbdplugin-rh2k8                     3/3     Running   0           9m52s
csi-rbdplugin-tbm56                     3/3     Running   0           9m52s
csi-rbdplugin-vq2lt                     3/3     Running   0           9m52s
rook-ceph-crashcollector-k8s-master03-79dbdbfddc-p7lcj  1/1     Running   0           8m35s
rook-ceph-crashcollector-k8s-work01-7fb8dfb6d6-w6f4r  1/1     Running   0           7m55s
rook-ceph-crashcollector-k8s-work02-6ccdf447-86m2k  1/1     Running   0           7m51s
rook-ceph-exporter-k8s-master03-765d886b4f-ls75w  1/1     Running   0           8m35s
rook-ceph-exporter-k8s-work01-6d444dfb4d-6qtsq  1/1     Running   0           7m50s
rook-ceph-exporter-k8s-work02-77f48b566f-zdngz  1/1     Running   0           7m46s
rook-ceph-mgr-a-59db64d547-qtcqj           3/3     Running   0           8m36s
rook-ceph-mgr-b-564896bcbf-p9wmn           3/3     Running   0           8m35s
rook-ceph-mon-a-6484dfb6c6-kc4xt           2/2     Running   0           9m30s
rook-ceph-mon-b-7d6cf98cc5-m62q7           2/2     Running   0           9m5s
rook-ceph-mon-c-9d8ff869f-skpmn            2/2     Running   0           8m52s
rook-ceph-operator-67944bdfcc-4h6st        1/1     Running   0           11m
rook-ceph-osd-0-85b8b478b9-s28b4           2/2     Running   0           7m55s
rook-ceph-osd-1-5d944c7777-vhm85           2/2     Running   0           7m53s
rook-ceph-osd-2-575df5bcbf-g6qsd           2/2     Running   0           7m51s
rook-ceph-osd-prepare-k8s-master03-6tvzj     0/1     Completed 0           8m12s
rook-ceph-osd-prepare-k8s-work01-cqbb7      0/1     Completed 0           8m12s
rook-ceph-osd-prepare-k8s-work02-42l2t      0/1     Completed 0           8m11s
rook-ceph-tools-7b75b967db-56plh           1/1     Running   0           37s
rook-discover-4wwbl                      1/1     Running   0           11m
rook-discover-cs2jr                      1/1     Running   0           11m
rook-discover-dftfs                      1/1     Running   0           11m
rook-discover-mrtcf                      1/1     Running   0           11m

```

如果第一次没有创建成功OSD，可清理重新尝试

```

1 # 停止当前集群的OSD服务（避免干扰）

```

```

2 kubectl -n rook-ceph delete --all pods -l app=rook-ceph-osd
3 #删除之前创建的crd
4 kubectl get crds | grep "ceph.rook.io" | awk '{print $1}' | xargs kubectl delete crd
5 kubectl get clusterroles | grep "rook" | awk '{print $1}' | xargs kubectl delete
  clusterrole
6 kubectl get clusterrolebindings | grep "rook" | awk '{print $1}' | xargs kubectl delete
  clusterrolebinding
7 kubectl get roles -n rook-ceph | awk '{print $1}' | xargs kubectl delete role -n rook-
  ceph
8 kubectl get rolebindings -n rook-ceph | awk '{print $1}' | xargs kubectl delete
  rolebinding -n rook-ceph
9 kubectl get serviceaccounts -n rook-ceph | awk '{print $1}' | xargs kubectl delete
  serviceaccount -n rook-ceph
10
11 # 在每个节点执行以下命令
12 sudo rm -rf /var/lib/rook/*
13 sudo wipefs -a /dev/nvme0n2
14 sudo dd if=/dev/zero of=/dev/nvme0n2 bs=1M count=1000
15
16 sudo reboot # 可选: 重启确保设备状态刷新

```

Note

需要注意的是，**osd-x** 的容器必须是存在的，且是正常的。如果上述 Pod 均正常，则认为集群安装成功。

更多配置: <https://rook.io/docs/rook/v1.6/ceph-cluster-crd.html>

2) 安装 ceph snapshot 控制器

snapshot 控制器的部署在集群安装时的 k8s-ha-install 项目中，需要切换到 1.20.x 分支：

```

1 cd /root/k8s-ha-install/
2 git checkout manual-installation-v1.20.x

```

创建 snapshot controller：

```

1 kubectl create -f /root/k8s-ha-install/snapshotter/ -n kube-system
2 kubectl get po -n kube-system -l app=snapshot-controller
3
4
5 [root@K8S-Master01 examples]# kubectl get po -n kube-system -l app=snapshot-controller
6 NAME                                READY   STATUS    RESTARTS   AGE
7 snapshot-controller-0               1/1     Running   0           6m16s

```


3) 安装 ceph 客户端工具

```
1 cd /root/rook/deploy/examples/
2 kubectl create -f /root/rook/deploy/examples/toolbox.yaml -n rook-ceph
3 kubectl get po -n rook-ceph -l app=rook-ceph-tools
4
5
6 [root@k8s-Master01 examples]# kubectl get po -n rook-ceph -l app=rook-ceph-tools
7 NAME                                READY   STATUS    RESTARTS   AGE
8 rook-ceph-tools-7b75b967db-1s9t4   1/1     Running   0           6m51s
```

容器启动后，即可进入容器内部执行服务命令

```
1 [root@k8s-master01 examples]# kubectl -n rook-ceph exec deploy/rook-ceph-tools -- ceph
  osd stat
2 3 osds: 3 up (since 7m), 3 in (since 7m); epoch: e16
3 [root@k8s-master01 examples]# kubectl -n rook-ceph exec deploy/rook-ceph-tools -- ceph
  -s
4   cluster:
5     id:      dd96e25d-c610-4fc2-a0aa-41bf4cd9d750
6     health: HEALTH_OK
7
8   services:
9     mon: 3 daemons, quorum a,b,c (age 5m)
10    mgr: b(active, since 6m), standbys: a
11    osd: 3 osds: 3 up (since 7m), 3 in (since 7m)
12
13   data:
14     pools: 1 pools, 1 pgs
15     objects: 2 objects, 449 KiB
16     usage: 82 MiB used, 30 GiB / 30 GiB avail
17     pgs: 1 active+clean
18
19 如果觉得执行命令过长,可以设置别名
```

具体文档: <https://rook.io/docs/rook/v1.6/ceph-csi-snapshot.html>

4) 安装Ceph dashboard

默认情况下，ceph dashboard是打开的，可以创建一个nodePort类型的Service暴露服务（新版本该文件默认存在可以直接创建）：

有存在以下四个类型的 SVC 文件，如果dashboard 之前设置的是ssl: true 那么根据实际情况去进行创建

```
1 [root@k8s-master01 examples]# ls dashboard-*
2 dashboard-external-https.yaml dashboard-external-http.yaml dashboard-ingress-
  https.yaml dashboard-loadbalancer.yaml
```

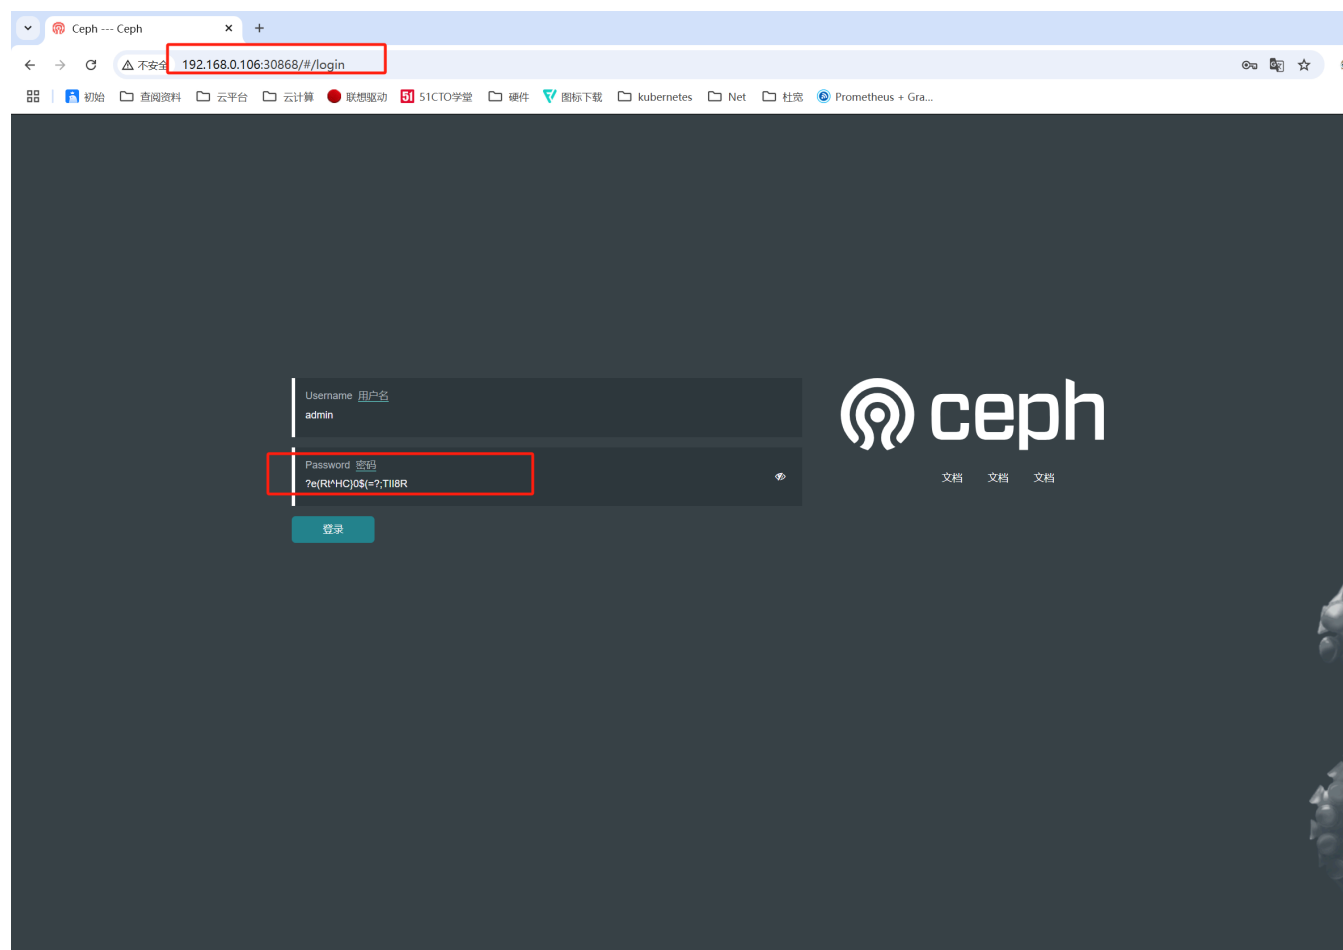
创建

```

1 kubectl create -f dashboard-external-http.yaml
2
3 [root@k8s-master01 examples]# kubectl get svc -n rook-ceph rook-ceph-mgr-dashboard-external-https
4
5 NAME                                TYPE        CLUSTER-IP    EXTERNAL-IP
6 PORT(S)          AGE
7 rook-ceph-mgr-dashboard-external-https  NodePort    10.96.211.249 <none>
8 8443:31983/TCP   80s
9
10 #获取登陆密码
11 kubectl -n rook-ceph get secret rook-ceph-dashboard-password -o jsonpath="{['data']['password']}" | base64 --decode && echo
12 ?e(Rt^HC}0$(=?;TII8R

```

WEB 访问



ceph 块存储的使用

块存储一般用于一个 Pod 挂载一块存储使用，相当于一个服务器新挂了一个盘，只给一个应用使用。

参考文档: <https://rook.io/docs/rook/v1.6/ceph-block.html>

1) 创建 StorageClass 和 ceph 的存储池

```
1 cd /root/rook/deploy/examples
2 kubectl create -f csi/rbd/storageclass.yaml -n rook-ceph
3
4
5 [root@k8s-master01 examples]# kubectl get cephblockpool -n rook-ceph
6 NAME          PHASE    TYPE          FAILUREDOMAIN  AGE
7 replicapool   Ready    Replicated    host           2m7s
8 [root@k8s-master01 examples]# kubectl get sc
9 NAME          PROVISIONER          RECLAIMPOLICY  VOLUMEBINDINGMODE
10 ALLOWVOLUMEEXPANSION  AGE
11 rook-ceph-block  rook-ceph.rbd.csi.ceph.com  Delete          Immediate          true
12 2m15s
13
14 # 查看ceph提供的存储驱动
15 kubectl get csidriver
```

```
[root@k8s-master01 examples]# kubectl get csidriver
NAME          ATTACHREQUIRED  PODINFOONMOUNT  STORAGECAPACITY  TOKENREQUESTS  REQUIRESREUBLISH  MODES          AGE
rook-ceph.cephfs.csi.ceph.com  true            false            false             <unset>         false             Persistent     13h
rook-ceph.rbd.csi.ceph.com     true            false            false             <unset>         false             Persistent     13h
```

Yaml 配置文件解释

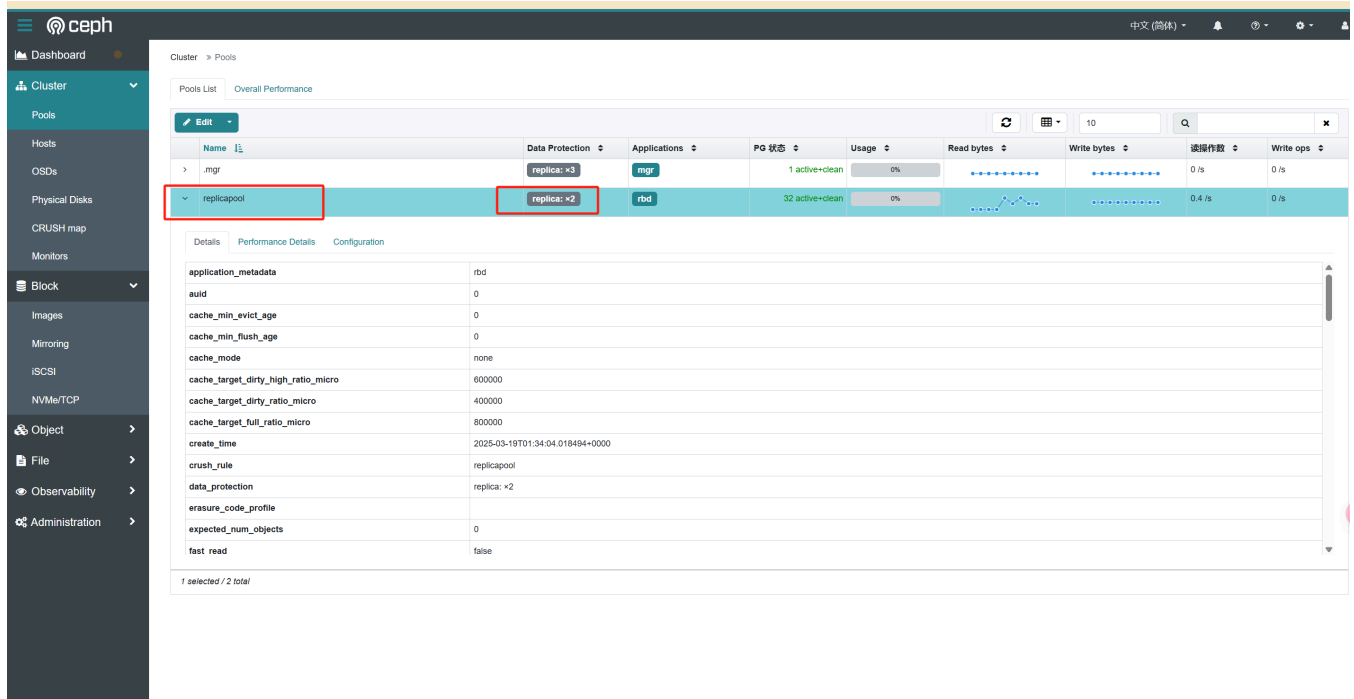
```
1 apiVersion: ceph.rook.io/v1
2 kind: CephBlockPool
3 metadata:
4   name: replicapool          # 存储池名称
5   namespace: rook-ceph      # 所属命名空间（必须与CephCluster一致）
6 spec:
7   failureDomain: host        # 数据副本分布策略（故障域级别）其他可选值：osd（不同OSD磁盘）、rack（不同机架）。
8   replicated:
9     size: 3                  # 数据副本数量,每个数据块保存 3个副本（即同一份数据在集群中有3份拷贝）。
10   requireSafeReplicaSize: true # 强制副本数必须满足最小安全要求
11 ---
12 apiVersion: storage.k8s.io/v1
13 kind: StorageClass
14 metadata:
15   name: rook-ceph-block      # 存储类名称（创建PVC时需指定）
16 provisioner: rook-ceph.rbd.csi.ceph.com # CSI驱动名称
17 parameters:
18   clusterID: rook-ceph      # Ceph集群ID（必须与CephCluster名称一致）
19   pool: replicapool         # 使用的Ceph存储池（即上述定义的replicapool）
20   imageFormat: "2"          # RBD镜像格式版本,使用 第2版RBD格式（支持更多功能如动态调整大小）。
21   imageFeatures: layering    # RBD镜像支持的特性,支持 分层克隆（用于快照和克隆功能）。
22   csi.storage.k8s.io/provisioner-secret-name: rook-csi-rbd-provisioner # 供应者密钥名称
23   csi.storage.k8s.io/provisioner-secret-namespace: rook-ceph # 密钥所在命名空间
24   csi.storage.k8s.io/controller-expand-secret-name: rook-csi-rbd-provisioner # 扩容存储卷时使用的认证信息名称
25   csi.storage.k8s.io/controller-expand-secret-namespace: rook-ceph # 密钥命名空间
26   csi.storage.k8s.io/node-stage-secret-name: rook-csi-rbd-node # 节点挂载密钥名称
```

```

27 | csi.storage.k8s.io/node-stage-secret-namespace: rook-ceph # 密钥命名空间
28 | csi.storage.k8s.io/fstype: ext4 # 文件系统类型,存储卷格式化时使用 ext4文件系统（可选: xfs、
    | ext3等）。
29 | allowVolumeExpansion: true # 允许通过Kubernetes动态扩展存储卷大小。
30 | reclaimPolicy: Delete # 回收策略,删除PVC时自动删除底层RBD镜像（可选: Retain 保留数据）

```

从Dashboard上可以看到刚刚所创建的Pool，并且副本数量是我们设置的 size: 2



2) 挂载测试

创建一个 MySQL 服务

```

1 | cd /root/rook/deploy/examples
2 | kubectl create -f mysql.yaml
3 |
4 | kubectl get pvc,pv,po

```

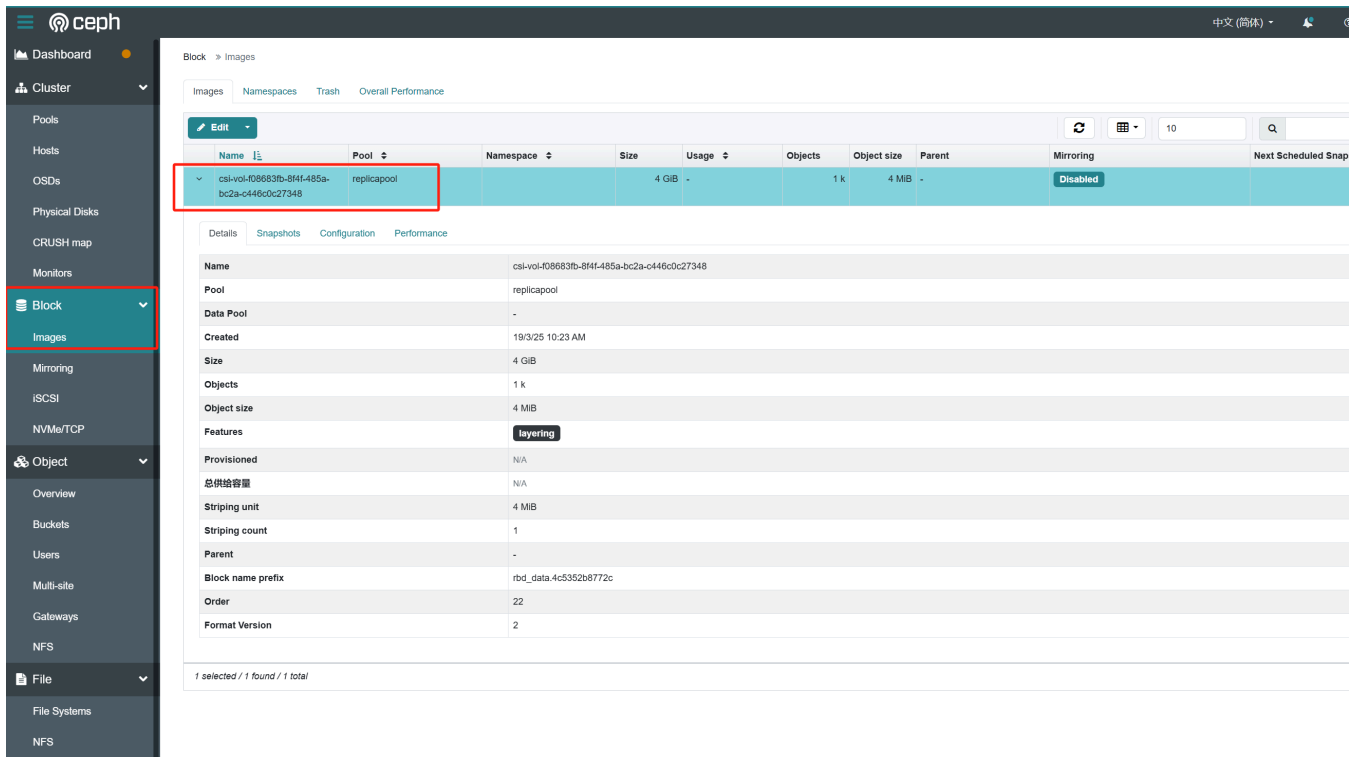
```

[root@k8s-master01 examples]# kubectl get pvc,pv,po
NAME                                STATUS    VOLUME                                     CAPACITY   ACCESS MODES   STORAGECLASS   VOLUMEATTRIBUTESCLASS   AGE
persistentvolumeclaim/mysql-pv-claim  Bound    pvc-52989a1f-ef09-4813-8623-514986a1f750  4Gi        RWO            rook-ceph-block <unset>              20m

NAME                                CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM                STORAGECLASS   VOLUMEATTRIBUTESCLASS   REASON   AGE
persistentvolume/pvc-52989a1f-ef09-4813-8623-514986a1f750  4Gi        RWO            Delete           Bound    default/mysql-pv-claim  rook-ceph-block <unset>              20m

NAME                                READY    STATUS    RESTARTS   AGE
pod/wordpress-mysql-cc5fd5cd9-rbsbl  1/1      Running   0           20m
[root@k8s-master01 examples]#

```



Yaml 配置文件解释

Note

Volume 通过主要配置参数: `claimName: mysql-pv-claim` 指定PVC

PVC 通过主要配置参数: `storageClassName: rook-ceph-block` 指定SC

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: wordpress-mysql      # 服务名称
5    labels:
6      app: wordpress           # 服务标签（用于关联应用）
7  spec:
8    ports:
9      - port: 3306              # 暴露的端口号（MySQL默认端口）
10   selector:
11     app: wordpress            # 选择器：匹配Pod的标签
12     tier: mysql
13   clusterIP: None             # 使用Headless Service（无集群IP）
14
15  ---
16  apiVersion: v1
17  kind: PersistentVolumeClaim
18  metadata:
19    name: mysql-pv-claim        # PVC名称
20    labels:
21      app: wordpress
22  spec:
23    storageClassName: rook-ceph-block # 指定使用的存储类（需提前创建）
24    accessModes:
25      - ReadWriteOnce           # 访问模式：单节点读写
```

```

26   resources:
27     requests:
28       storage: 5Gi          # 请求5Gi存储空间
29   ---
30   apiVersion: apps/v1
31   kind: Deployment
32   metadata:
33     name: wordpress-mysql   # 部署名称
34     labels:
35       app: wordpress
36       tier: mysql
37   spec:
38     selector:
39       matchLabels:          # 选择器：匹配Pod标签
40         app: wordpress
41         tier: mysql
42     strategy:
43       type: Recreate        # 更新策略：先终止旧Pod再创建新Pod（防止数据冲突）
44     template:
45       metadata:
46         labels:
47           app: wordpress
48           tier: mysql
49       spec:
50         containers:
51           - image: mysql:5.6  # 使用MySQL 5.6镜像
52             name: mysql
53             env:
54               - name: MYSQL_ROOT_PASSWORD # 设置MySQL root密码
55                 value: changeme          # 实际环境应使用Secret管理
56             ports:
57               - containerPort: 3306      # 容器暴露端口
58                 name: mysql
59             volumeMounts:
60               - name: mysql-persistent-storage
61                 mountPath: /var/lib/mysql # MySQL数据存储路径
62         volumes:
63           - name: mysql-persistent-storage
64             persistentVolumeClaim:
65               claimName: mysql-pv-claim   # 绑定到上述PVC

```

pvc 会连接刚才创建的 `storageClass`，动态的创建 `pv`，然后连接到 ceph 创建对应的存储之后创建。`pvc` 只需要指定 `storageClassName` 为刚才创建的 `StorageClass` 名称即可连接到 rook 的ceph。如果是 statefulset，只需要将 `volumeTemplateClaim` 里面的 `claim` 名称改为 `StorageClass` 名称即可动态创建 Pod

完整流程说明

1.存储分配：

- 当PVC `mysql-pv-claim` 被创建时，Rook-Ceph会根据 `rook-ceph-block` 存储类动态创建PV（RBD镜像）。
- PVC与PV绑定后，MySQL Pod才能挂载存储卷。

2.服务访问：

- 其他Pod（如WordPress）可通过DNS名称 `wordpress-mysql` 访问MySQL服务。
- Headless Service的DNS记录直接指向MySQL Pod IP，适用于需要直接访问Pod的场景。

3.数据持久化:

- MySQL数据保存在 `/var/lib/mysql` 目录，底层由Ceph RBD提供高可用存储。
- 即使Pod重启或迁移，数据不会丢失。

卸载pv

```
1 [root@k8s-master01 examples]# kubectl delete -f mysql.yaml
2 service "wordpress-mysql" deleted
3 persistentvolumeclaim "mysql-pv-claim" deleted
4 deployment.apps "wordpress-mysql" deleted
5
6
7 [root@k8s-master01 examples]# kubectl get pvc,pv,po
8 No resources found
```

3) StatefulSet 动态存储

普通情况 PVC通过 StorageClass 一次只能创一个PV，但是 StatefulSet 有状态的服务不适用于共享存储数据，否则可能会导致服务异常，所以每个po需要一个pv。可以通过volumeClaimTemplates：来实现

```
1 #部署
2 vim sts-sc.yaml
3 kubectl create -f sts-sc.yaml
4
5 kubectl get pvc,pv,po
6
7
8 #卸载
9 kubectl delete -f sts-sc.yaml
10 kubectl delete pvc www-web-0,www-web-1,www-web-2
```

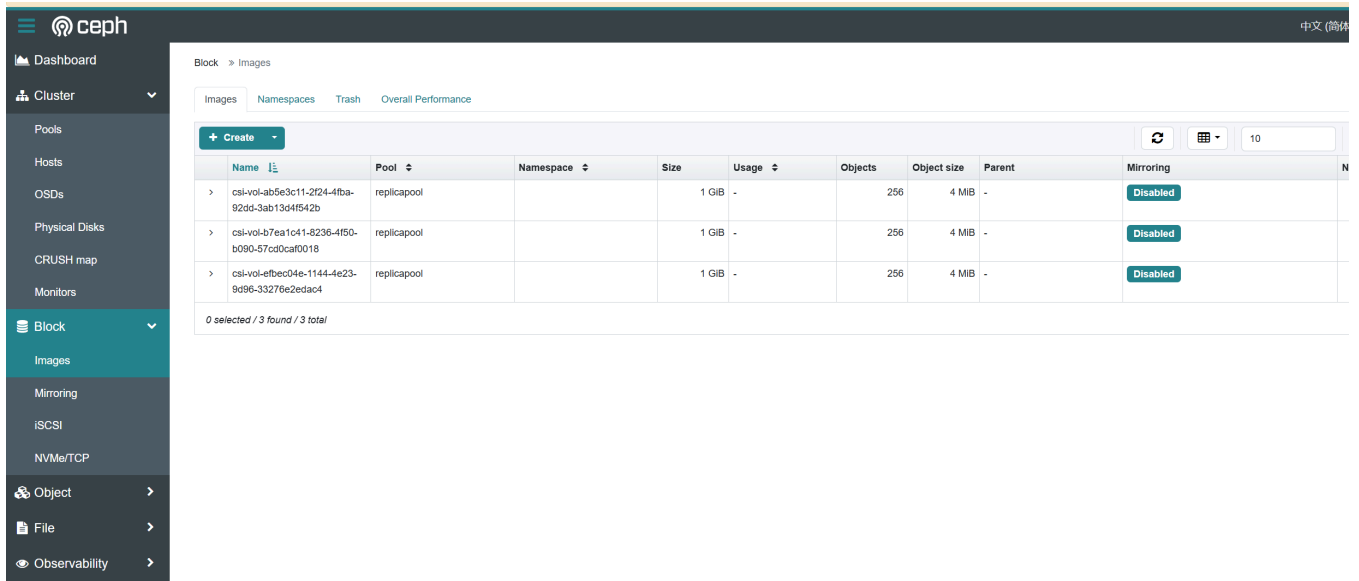
```
[root@k8s-master01 examples]# kubectl get pvc,pv,po
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	VOLUMEATTRIBUTESCLASS	AGE
persistentvolumeclaim/www-web-0	Bound	pvc-6fc09f4a-adee-40d3-9490-6389dab672c8	1Gi	RWO	rook-ceph-block	<unset>	52m
persistentvolumeclaim/www-web-1	Bound	pvc-8bd40528-b945-42c8-9aeb-8df10e42548f	1Gi	RWO	rook-ceph-block	<unset>	20m
persistentvolumeclaim/www-web-2	Bound	pvc-56d3b62b-a2ba-4d4a-aabf-fe9c58cffe71	1Gi	RWO	rook-ceph-block	<unset>	20m

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM	STORAGECLASS	VOLUMEATTRIBUTESCLASS	REASON	AGE
persistentvolume/pvc-56d3b62b-a2ba-4d4a-aabf-fe9c58cffe71	1Gi	RWO	Delete	Bound	default/www-web-2	rook-ceph-block	<unset>		20m
persistentvolume/pvc-6fc09f4a-adee-40d3-9490-6389dab672c8	1Gi	RWO	Delete	Bound	default/www-web-0	rook-ceph-block	<unset>		52m
persistentvolume/pvc-8bd40528-b945-42c8-9aeb-8df10e42548f	1Gi	RWO	Delete	Bound	default/www-web-1	rook-ceph-block	<unset>		20m

NAME	READY	STATUS	RESTARTS	AGE
pod/web-0	1/1	Running	0	18m
pod/web-1	1/1	Running	0	17m
pod/web-2	1/1	Running	0	17m

```
[root@k8s-master01 examples]#
```



```
1  # 定义一个 Service 资源
2  apiVersion: v1
3  kind: Service
4  metadata:
5    name: nginx                # 服务名称
6    labels:
7      app: nginx                # 服务标签（用于关联应用）
8  spec:
9    ports:
10     - port: 80                # 服务暴露端口
11       name: web                # 端口名称标识
12     clusterIP: None           # 使用Headless Service（无集群IP）
13     selector:
14       app: nginx              # 选择器：匹配Pod的标签
15
16  ---
17  # 定义一个 StatefulSet 资源
18  apiVersion: apps/v1
19  kind: StatefulSet
20  metadata:
21    name: web                  # StatefulSet名称
22  spec:
23    selector:
24      matchLabels:
25        app: nginx             # 必须与Pod模板中的标签一致
26    serviceName: "nginx"        # 关联的Headless Service名称
27    replicas: 3                 # 副本数量（创建3个Pod: web-0, web-1, web-2）
28    minReadySeconds: 10         # 新Pod就绪后等待10秒才视为可用
29    template:
30      metadata:
31        labels:
32          app: nginx            # Pod标签（必须与selector.matchLabels一致）
33      spec:
34        terminationGracePeriodSeconds: 10 # 删除Pod时的优雅终止等待时间
35        containers:
36          - name: nginx
```



```

37     image: m.daocloud.io/docker.io/library/nginx:latest # 使用Nginx镜像
38     ports:
39     - containerPort: 80 # 容器监听端口
40       name: web # 端口名称标识
41     volumeMounts:
42     - name: www # 挂载的卷名称（与volumeClaimTemplates匹配）
43       mountPath: /usr/share/nginx/html # 挂载路径（Nginx默认静态文件目录）
44     volumeClaimTemplates: # 动态创建PVC的模板（每个Pod自动生成独立PVC）
45     - metadata:
46       name: www # PVC名称模板（最终名称：www-web-0，www-web-1等）
47     spec:
48       accessModes: [ "ReadWriteOnce" ] # 访问模式：单节点读写
49       storageClassName: "rook-ceph-block" # 使用的存储类（需提前创建）
50       resources:
51       requests:
52         storage: 1Gi # 每个PVC请求1Gi存储空间

```

共享文件系统的使用

共享文件系统一般用于多个 Pod 共享一个存储

官方文档: <https://rook.io/docs/rook/v1.6/ceph-filestorage.html>

1) 创建共享类型的文件系统和 StorageClass

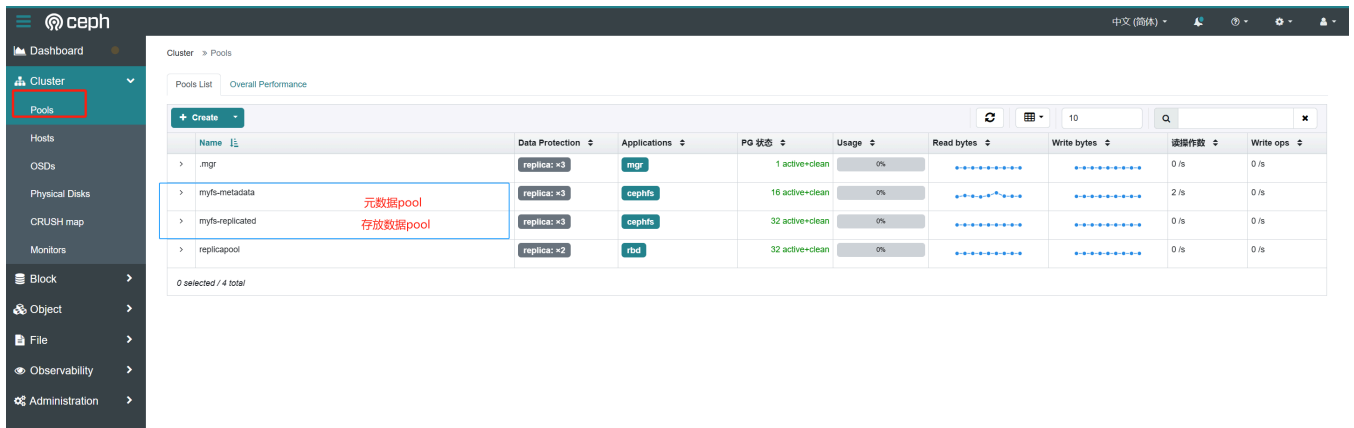
📌 Note

文件存储会有创建MDS 用来存放元数据，多个 MDS 节点共享元数据，通过动态子树分区实现负载均衡

```

1  cd /root/rook/deploy/examples
2  kubectl create -f filesystem.yaml
3  kubectl -n rook-ceph get pod -l app=rook-ceph-mds
4
5
6  kubectl create -f csi/cephfs/storageclass.yaml
7
8  [root@k8s-master01 examples]# kubectl get sc
9  NAME                                PROVISIONER                                RECLAIMPOLICY    VOLUMEBINDINGMODE
10  rook-ceph-block                    rook-ceph.rbd.csi.ceph.com                Delete           Immediate
11  rook-cephfs                        rook-ceph.cephfs.csi.ceph.com             Delete           Immediate

```



```
[root@k8s-master01 examples]# kubectl -n rook-ceph get pod -l app=rook-ceph-mds
NAME                                READY   STATUS    RESTARTS   AGE
rook-ceph-mds-myfs-a-6977cb87f-7p8h5 2/2     Running   0           11m
rook-ceph-mds-myfs-b-75559576b7-mnd6f 2/2     Running   0           11m
[root@k8s-master01 examples]#
```

2) nginx 挂载测试

```
1 kubectl create -f nginx.yaml
2 kubectl get po -l app=nginx
3
4 kubectl get pvc,pv,po
5
6 vim nginx.yaml
```

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: nginx                # 服务名称
5   labels:
6     app: nginx               # 服务标签（用于关联应用）
7 spec:
8   ports:
9     - port: 80               # 服务暴露端口
10     name: web                # 端口名称标识
11   selector:
12     app: nginx               # 选择器：匹配Pod的标签,将流量路由到标签为 app: nginx 的Pod。
13   type: ClusterIP           # 服务类型（默认ClusterIP，仅集群内访问）
14 ---
15 kind: PersistentVolumeClaim
16 apiVersion: v1
17 metadata:
18   name: nginx-share-pvc      # PVC名称
19 spec:
20   storageClassName: rook-cephfs # 使用的存储类（需提前创建CephFS存储类）
21   accessModes:
22     - ReadWriteMany          # 访问模式：多节点读写,允许多个Pod同时读写同一个存储卷（适合共享存储场景）
23   resources:
24     requests:
25       storage: 1Gi           # 请求1Gi存储空间
26 ---
```

```

27 apiVersion: apps/v1
28 kind: Deployment
29 metadata:
30   name: web # 部署名称
31 spec:
32   selector:
33     matchLabels:
34       app: nginx # 选择器：匹配Pod标签
35   replicas: 3 # 副本数量（部署3个Pod）
36   template:
37     metadata:
38       labels:
39         app: nginx # Pod标签（必须与selector.matchLabels一致）
40   spec:
41     containers:
42     - name: nginx
43       image: m.daocloud.io/docker.io/library/nginx:latest # 使用Nginx镜像
44       imagePullPolicy: IfNotPresent # 镜像拉取策略（本地存在则不拉取）
45       ports:
46       - containerPort: 80 # 容器监听端口
47         name: web
48       volumeMounts:
49       - name: www # 挂载的卷名称
50         mountPath: /usr/share/nginx/html # Nginx静态文件目录
51     volumes:
52     - name: www # 卷名称（与volumeMounts匹配）
53       persistentVolumeClaim:
54         claimName: nginx-share-pvc # 绑定到上述PVC

```

```

[root@k8s-master01 examples]# kubectl get po -l app=nginx -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE      NOMINATED NODE   READINESS GATES
web-86c6654564-6n8pz  1/1     Running   0          79s   172.16.122.150 k8s-master02  <none>          <none>
web-86c6654564-84gm5  1/1     Running   0          79s   172.16.236.51  k8s-work01   <none>          <none>
web-86c6654564-dzr57  1/1     Running   0          79s   172.16.178.169 k8s-work02   <none>          <none>

[root@k8s-master01 examples]# kubectl get po,pv,pvc
NAME          READY   STATUS    RESTARTS   AGE
pod/web-86c6654564-6n8pz  1/1     Running   0          81s
pod/web-86c6654564-84gm5  1/1     Running   0          81s
pod/web-86c6654564-dzr57  1/1     Running   0          81s

NAME          CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM                STORAGECLASS   VOLUMEATTRIBUTESCLASS   REASON   AGE
persistentvolume/pvc-8b9641d7-2a59-4323-bb30-719762206ed2  1Gi       RWX            Delete           Bound    default/nginx-share-pvc  rook-cephfs   <unset>                  76s

NAME          STATUS   VOLUME          CAPACITY   ACCESS MODES   STORAGECLASS   VOLUMEATTRIBUTESCLASS   AGE
persistentvolumeclaim/nginx-share-pvc  Bound    pvc-8b9641d7-2a59-4323-bb30-719762206ed2  1Gi       RWX            rook-cephfs   <unset>                  82s

```

3) 验证效果

```

[root@k8s-master01 examples]# kubectl get po
NAME          READY   STATUS    RESTARTS   AGE
web-86c6654564-6n8pz  1/1     Running   0          20m
web-86c6654564-84gm5  1/1     Running   0          20m
web-86c6654564-dzr57  1/1     Running   0          20m

[root@k8s-master01 examples]# kubectl exec -it po web-86c6654564-6n8pz -- mkdir /usr/share/nginx/html/1
Error from server (NotFound): pods "po" not found
[root@k8s-master01 examples]# kubectl exec -it web-86c6654564-6n8pz -- mkdir /usr/share/nginx/html/1
[root@k8s-master01 examples]# kubectl exec -it web-86c6654564-6n8pz -- ls /usr/share/nginx/html/1
1
[root@k8s-master01 examples]# kubectl exec -it web-86c6654564-84gm5 -- ls /usr/share/nginx/html/1
1

```

PVC 扩容

⚠ Caution

文件共享类型的 PVC 扩容需要 k8s 1.15+
块存储类型的 PVC 扩容需要 k8s 1.16+

在 Kubernetes 中，PVC 的缩容功能是受限制的。大多数存储提供商不支持 PVC 的缩容操作。即使某些存储提供商支持缩容，也需要手动干预，并且可能会导致数据丢失。因此，Kubernetes 默认不允许 PVC 的缩容操作。

1) 扩容文件共享型 PVC

前置条件：需确保 SC 开启了参数 allowVolumeExpansion: true

```
1 [root@k8s-master01 examples]# awk '/allowVolumeExpansion/' csi/cephfs/storageclass.yaml
2 allowVolumeExpansion: true
```

修改扩容

```
1 kubectl get pvc,pvc
2 kubectl edit pvc cephfs-pvc -n kube-system
3
4
5 # 查看修改过后的容量大小
6 kubectl get pvc,pvc
```

```
[root@k8s-master01 examples]# kubectl get pvc,pv
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS VOLUMEATTRIBUTESCLASS AGE
persistentvolumeclaim/nginx-share-pvc Bound pvc-8b9641d7-2a59-4323-bb30-719762206ed2 1Gi RWX rook-cephfs <unset> 5h3m 修改之前容量为1G

NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM STORAGECLASS VOLUMEATTRIBUTESCLASS REASON AGE
persistentvolume/pvc-8b9641d7-2a59-4323-bb30-719762206ed2 1Gi RWX Delete Bound default/nginx-share-pvc rook-cephfs <unset> 5h3m
[root@k8s-master01 examples]# kubectl edit pvc nginx-share-pvc
persistentvolumeclaim/nginx-share-pvc edited
[root@k8s-master01 examples]# kubectl get pvc,pv
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS VOLUMEATTRIBUTESCLASS AGE
persistentvolumeclaim/nginx-share-pvc Bound pvc-8b9641d7-2a59-4323-bb30-719762206ed2 2Gi RWX rook-cephfs <unset> 5h4m 修改之后容量为自定义的2G

NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM STORAGECLASS VOLUMEATTRIBUTESCLASS REASON AGE
persistentvolume/pvc-8b9641d7-2a59-4323-bb30-719762206ed2 2Gi RWX Delete Bound default/nginx-share-pvc rook-cephfs <unset> 5h4m
[root@k8s-master01 examples]#
```

2) 扩容块存储

与文件系统扩容一样，先确认sc有没有开启允许动态扩容,然后直接 edit pv进行容量修改

```
1 kubectl get pvc,pv,pv
2 kubectl edit pvc mysql-pv-claim
```

```
[root@k8s-master01 examples]# kubectl get pvc,pv,pv
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS VOLUMEATTRIBUTESCLASS AGE
persistentvolumeclaim/mysql-pv-claim Bound pvc-903cbe1d-8c37-4556-bbdf-65f235a31dcc 4Gi RWO rook-ceph-block <unset> 62s

NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM STORAGECLASS VOLUMEATTRIBUTESCLASS REASON AGE
persistentvolume/pvc-903cbe1d-8c37-4556-bbdf-65f235a31dcc 4Gi RWO Delete Bound default/mysql-pv-claim rook-ceph-block <unset> 61s 修改前

NAME pod/wordpress-mysql-cc5fd5cd9-7qd7c READY STATUS RESTARTS AGE
1/1 Running 0 62s
[root@k8s-master01 examples]# kubectl edit pvc mysql-pv-claim
persistentvolumeclaim/mysql-pv-claim edited
[root@k8s-master01 examples]# kubectl get pvc,pv,pv
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS VOLUMEATTRIBUTESCLASS AGE
persistentvolumeclaim/mysql-pv-claim Bound pvc-903cbe1d-8c37-4556-bbdf-65f235a31dcc 5Gi RWO rook-ceph-block <unset> 4m18s

NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM STORAGECLASS VOLUMEATTRIBUTESCLASS REASON AGE
persistentvolume/pvc-903cbe1d-8c37-4556-bbdf-65f235a31dcc 5Gi RWO Delete Bound default/mysql-pv-claim rook-ceph-block <unset> 4m17s 修改后，pvc显示容量状态需要等待一会儿刷新

NAME pod/wordpress-mysql-cc5fd5cd9-7qd7c READY STATUS RESTARTS AGE
1/1 Running 0 4m18s
```

```

# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will be
# reopened with the relevant failures.
#
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  annotations:
    pv.kubernetes.io/bind-completed: "yes"
    pv.kubernetes.io/bound-by-controller: "yes"
    volume.beta.kubernetes.io/storage-provisioner: rook-ceph.rbd.csi.ceph.com
    volume.kubernetes.io/storage-provisioner: rook-ceph.rbd.csi.ceph.com
  creationTimestamp: "2025-03-20T07:31:01Z"
  finalizers:
    - kubernetes.io/pvc-protection
  labels:
    app: wordpress
    name: mysql-pv-claim
    namespace: default
    resourceVersion: "462738"
    uid: 903cbe1d-8c37-4556-bbdf-65f235a31dcc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 6Gi
  storageClassName: rook-ceph-block
  volumeMode: Filesystem
  volumeName: pvc-903cbe1d-8c37-4556-bbdf-65f235a31dcc
status:
  accessModes:
    - ReadWriteOnce
  capacity:
    storage: 4Gi
  phase: Bound
~
~
~

```

修改容量大小

PVC 快照

1) 创建 snapshotClass

```

1 kubectl create -f csi/rbd/snapshotclass.yaml
2
3 [root@k8s-master01 examples]# cat csi/rbd/snapshotclass.yaml
4 ---
5 apiVersion: snapshot.storage.k8s.io/v1
6 kind: VolumeSnapshotClass
7 metadata:
8   name: csi-rbdplugin-snapclass
9   driver: rook-ceph.rbd.csi.ceph.com # csi-provisioner-name
10  parameters:
11    # Specify a string that identifies your cluster. Ceph CSI supports any
12    # unique string. When Ceph CSI is deployed by Rook use the Rook namespace,
13    # for example "rook-ceph".
14    clusterID: rook-ceph # namespace:cluster
15    csi.storage.k8s.io/snapshotter-secret-name: rook-csi-rbd-provisioner
16    csi.storage.k8s.io/snapshotter-secret-namespace: rook-ceph # namespace:cluster
17  deletionPolicy: Delete

```

2) 创建快照

创建一些数据模拟环境

```

1 kubectl exec -it wordpress-mysql-cc5fd5cd9-7qd7c -- mkdir /var/lib/mysql/demo{1..3}
2 kubectl get po,pvc,volumeSnapshotClass

```

```

3
4
5 [root@k8s-master01 examples]# vim snapshot.yaml
6 ---
7 apiVersion: snapshot.storage.k8s.io/v1
8 kind: VolumeSnapshot
9 metadata:
10   name: rbd-pvc-snapshot
11 spec:
12   volumeSnapshotClassName: csi-rbdplugin-snapclass
13   source:
14     persistentVolumeClaimName: rbd-pvc
15
16
17 kubectl create -f snapshot.yaml

```

```

[root@k8s-master01 examples]# kubectl get po,pvc,volumeSnapshotClass
NAME                                READY   STATUS    RESTARTS   AGE
pod/wordpress-mysql-cc5fd5cd9-rx67c 1/1     Running   0           26m

NAME                                STATUS    VOLUME                                     CAPACITY   ACCESS MODES   STORAGECLASS   VOLUMEATTRIBUTESCLASS   AGE
persistentvolumeclaim/mysql-pv-claim Bound     pvc-09c59b1e-089e-472b-86c0-d3db5aaddde7 4Gi         RWO            rook-ceph-block <unset>                 26m

NAME                                VOLUMESnapshotClass.snapshot.storage.k8s.io/csi-rbdplugin-snapclass DRIVER rook-ceph.rbd.csi.ceph.com DELETIONPOLICY Delete AGE 4h12m
[root@k8s-master01 examples]# cat snapshot.yaml
---
apiVersion: snapshot.storage.k8s.io/v1
kind: VolumeSnapshot
metadata:
  name: rbd-pvc-snapshot
spec:
  volumeSnapshotClassName: csi-rbdplugin-snapclass
  source:
    persistentVolumeClaimName: mysql-pv-claim
[root@k8s-master01 examples]# kubectl get sc
NAME            PROVISIONER             RECLAIMPOLICY   VOLUMEINDELMODE   ALLOWVOLUMEEXPANSION   AGE
rook-ceph-block rook-ceph.rbd.csi.ceph.com Delete           Immediate          true                   36h
rook-cephfs     rook-ceph.cephfs.csi.ceph.com Delete           Immediate          true                   11h
[root@k8s-master01 examples]# kubectl create -f snapshot.yaml
volumesnapshot.snapshot.storage.k8s.io/rbd-pvc-snapshot created
[root@k8s-master01 examples]# kubectl get -f snapshot.yaml
NAME                                READYTOUSE   SOURCEPVC   SOURCESNAPSHOTCONTENT   RESTORESIZE   SNAPSHOTCLASS   SNAPSHOTCONTENT   CREATIONTIME   AGE
rbd-pvc-snapshot true          mysql-pv-claim snapcontent-0f82a7cd-e84c-409e-80e0-6c54cf047379 4Gi           csi-rbdplugin-snapclass snapcontent-0f82a7cd-e84c-409e-80e0-6c54cf047379 3s             5s
[root@k8s-master01 examples]#

```

查看快照

```

1 [root@k8s-master01 examples]# kubectl get -f snapshot.yaml
2 NAME                                READYTOUSE   SOURCEPVC   SOURCESNAPSHOTCONTENT   RESTORESIZE
SNAPSHOTCLASS                                SNAPSHOTCONTENT
CREATIONTIME   AGE
3 rbd-pvc-snapshot true          mysql-pv-claim 4Gi
csi-rbdplugin-snapclass snapcontent-0f82a7cd-e84c-409e-80e0-6c54cf047379 4m7s
4m9s

```

3) 通过快照创建 PVC

如果想要创建一个具有某个数据的 PVC，可以从某个快照恢复

```

1 cat pvc-restore.yaml
2 apiVersion: v1 # API版本声明（Kubernetes核心API）
3 kind: PersistentVolumeClaim # 资源类型为持久卷声明（PVC）
4 metadata:
5   name: rbd-pvc-restore # PVC名称（用于标识该存储声明）
6 spec:
7   storageClassName: rook-ceph-block # 存储类名称（指向预先创建的rook-ceph-block存储类）
8   dataSource:
9     name: rbd-pvc-snapshot # 引用的VolumeSnapshot名称（需预先存在）
10    kind: volumeSnapshot # 资源类型为存储快照

```

```

11     apiGroup: snapshot.storage.k8s.io # 快照资源所属的API组
12     accessModes:
13       - ReadWriteOnce # 访问模式配置：单节点读写挂载模式
14     resources:
15       requests:
16         storage: 4Gi # 存储空间请求量（必须 >= 原始PVC容量）
17
18 kubectl create -f pvc-restore.yaml

```

```

[root@k8s-master01 examples]# kubectl create -f pvc-restore.yaml
persistentvolumeclaim/rbd-pvc-restore created
[root@k8s-master01 examples]# cat pvc-restore.yaml
apiVersion: v1 # API版本声明 (Kubernetes核心API)
kind: PersistentVolumeClaim # 资源类型为持久卷声明 (PVC)
metadata:
  name: rbd-pvc-restore # PVC名称 (用于标识该存储声明)
spec:
  storageClassName: rook-ceph-block # 存储类名称 (指向预先创建的rook-ceph-block存储类)
  dataSource:
    name: rbd-pvc-snapshot # 引用的VolumeSnapshot名称 (需预先存在)
    kind: VolumeSnapshot # 资源类型为存储快照
    apiGroup: snapshot.storage.k8s.io # 快照资源所属的API组
  accessModes:
    - ReadWriteOnce # 访问模式配置：单节点读写挂载模式
  resources:
    requests:
      storage: 4Gi # 存储空间请求量 (必须 >= 原始PVC容量)

```

通过快照创建的pvc

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	VOLUMEATTRIBUTESCLASS	AGE
mysql-pv-claim	Bound	pvc-09c59b1e-089e-472b-86c0-d3db5aaddde7	4Gi	RWO	rook-ceph-block	<unset>	58m
rbd-pvc-restore	Bound	pvc-b1ec2323-e0db-486c-95ea-7ad0c3746f1e	4Gi	RWO	rook-ceph-block	<unset>	18s

```

[root@k8s-master01 examples]#

```

4) 数据校验

创建pod绑定pvc验证校验数据

```

1 vim restore-check-snapshot-rbd.yaml
2 ---
3 apiVersion: apps/v1 # API版本声明 (适用于Deployment资源)
4 kind: Deployment # 资源类型为Deployment
5 metadata:
6   name: check-snapshot-restore # Deployment的名称
7 spec:
8   selector:
9     matchLabels:
10      app: check # 匹配Pod的标签，用于关联Deployment和Pod
11   strategy:
12     type: Recreate # 部署策略：先删除旧Pod，再创建新Pod
13   template:
14     # pod 的模板
15     metadata:
16       labels:
17         app: check # Pod的标签，与selector中的标签匹配
18     spec:
19       containers:
20         - image: m.daocloud.io/docker.io/library/nginx:latest # 使用的容器镜像
21           name: check # 容器的名称
22           command:
23             - sh
24             - -c
25             - sleep 36000 # 容器启动后执行的命令：休眠36000秒（10小时）
26           volumeMounts:
27             - name: check-mysql-persistent-storage # 挂载的卷名称
28               mountPath: /mnt # 卷挂载到容器内的路径
29       volumes:

```

```

29     - name: check-mysql-persistent-storage # 卷的名称
30     persistentVolumeClaim:
31       claimName: rbd-pvc-restore # 使用的PVC名称（从快照恢复的PVC）
32
33   kubectl create -f restore-check-snapshot-rbd.yaml
34   kubectl get -f restore-check-snapshot-rbd.yaml

```

```

[root@k8s-master01 examples]# kubectl exec -it check-snapshot-restore-78f558698-4rqrp -- ls /mnt
auto.cnf      ib_logfile1  lost+found    performance_schema
ib_logfile0  ibddatal     mysql         test
[root@k8s-master01 examples]# kubectl exec -it check-snapshot-restore-78f558698-4rqrp -- ls /mnt/mysql
columns_priv.MYD  func.MYD      help_keyword.MYD  innodb_index_stats.ibd  proc.MYD      servers.MYD      slave_worker_info.ibd  time_zone.frm      time_zone_transition.frm
columns_priv.frm  func.frm      help_keyword.frm  innodb_index_stats.frm  proc.frm      servers.frm      slave_worker_info.frm  time_zone.frm      time_zone_transition.frm
db.MYD            general_log.CSM  help_relation.MYD  innodb_table_stats.ibd  proc.frm      servers.frm      slave_worker_info.frm  time_zone.frm      time_zone_transition.frm
db.MYD            general_log.CSV  help_relation.MYD  ndb_binlog_index.MYD    proc_priv.MYD  slave_master_info.frm  tables_priv.MYD        time_zone.frm      time_zone_transition.frm
db.frm           general_log.frm  help_relation.frm  ndb_binlog_index.frm    proc_priv.frm  slave_master_info.frm  tables_priv.frm        time_zone.frm      time_zone_transition.frm
event.MYD         help_category.MYD  help_topic.MYD     ndb_binlog_index.frm    proc_priv.frm  slave_relay_log_info.frm  tables_priv.frm        time_zone.frm      time_zone_transition.frm
event.MYD         help_category.frm  help_topic.frm     plugin.MYD               proxies_priv.MYD  slave_relay_log_info.frm  tables_priv.frm        time_zone.frm      time_zone_transition.frm
event.frm         help_category.frm  help_topic.frm     plugin.MYD               proxies_priv.MYD  slave_relay_log_info.frm  tables_priv.frm        time_zone.frm      time_zone_transition.frm

```

5) 文件共享类型快照

操作步骤和块存储类型无区别，可以参考：

<https://rook.io/docs/rook/v1.6/ceph-csi-snapshot.html#cephfs-snapshots>

PVC 克隆

需要注意的是 pvc-clone.yaml 的 dataSource 的 name 是被克隆的 pvc 名称，在此是 mysql-pvclaim，storageClassName 为新建 pvc 的 storageClass 名称，storage 不能小于之前 pvc 的大小。

```

1  vim pvc-clone.yaml
2  ---
3  apiVersion: v1 # API版本声明（Kubernetes核心API）
4  kind: PersistentVolumeClaim # 资源类型为持久卷声明（PVC）
5  metadata:
6    name: rbd-pvc-clone # PVC的名称（用于标识该克隆PVC）
7  spec:
8    storageClassName: rook-ceph-block # 存储类名称（指向预先创建的rook-ceph-block存储类）
9    dataSource:
10     name: mysql-pv-claim # 数据源名称（引用的原始PVC名称）
11     kind: PersistentVolumeClaim # 数据源类型为PVC（表示从现有PVC克隆）
12   accessModes:
13     - ReadWriteOnce # 访问模式配置：单节点读写挂载模式
14   resources:
15     requests:
16       storage: 3Gi # 存储空间请求量（必须 >= 原始PVC容量）
17
18   kubectl create -f pvc-clone.yaml

```

测试数据清理

参考链接：<https://rook.io/docs/rook/v1.6/ceph-teardown.html>

如果 Rook 要继续使用，可以只清理创建的 `deploy`、`pod`、`pvc` 即可。

清理步骤：

1.清理挂载了 PVC 的 Pod 和 Deployment

- 删除所有挂载了 PVC 的 Pod、Deployment 或其他高级资源。

```
1 kubectl delete pod <pod-name>
2 kubectl delete deployment <deployment-name>
3
4 [root@k8s-master01 examples]# kubectl get deploy
5 NAME                                READY    UP-TO-DATE    AVAILABLE    AGE
6 check-snapshot-restore              1/1      1              1            92m
7 wordpress-mysql                     1/1      1              1            95m
8 [root@k8s-master01 examples]# kubectl delete deploy check-snapshot-restore
9 deployment.apps "check-snapshot-restore" deleted
10 deployment.apps "wordpress-mysql" deleted
11 [root@k8s-master01 examples]# kubectl get po
12 NAME                                READY    STATUS              RESTARTS    AGE
13 check-snapshot-restore-78f558698-4rgrp 1/1      Terminating        0            89m
14 [root@k8s-master01 examples]# kubectl delete po check-snapshot-restore-78f558698-4rgrp
```

2.清理 PVC

- 删除所有通过 Ceph StorageClass 创建的 PVC。
- 检查 PV 是否被自动清理。

```
1 kubectl delete pvc <pvc-name>
2 kubectl get pv # 确认 PV 是否已清理
3
4
5
6 [root@k8s-master01 examples]# kubectl get pvc
7 NAME                                STATUS    VOLUME                                     CAPACITY
8 mysql-pv-claim                      Bound     pvc-dbea43d6-3991-4434-86f7-28707b55f2cb 4Gi
9 rbd-pvc-restore                      Bound     pvc-b1ec2323-e0db-486c-95ea-7ad0c3746f1e 4Gi
10 persistentvolumeclaim "mysql-pv-claim" deleted
11 persistentvolumeclaim "rbd-pvc-restore" deleted
12 [root@k8s-master01 examples]# kubectl get pv
13 No resources found
```

3.清理快照

- 删除所有 VolumeSnapshot 资源。

```

1 kubectl delete volumesnapshot <snapshot-name>
2
3
4 [root@k8s-master01 examples]# kubectl get volumesnapshot
5 NAME                                READYTOUSE    SOURCEPVC    SOURCESNAPSHOTCONTENT
6 RESTORESIZE    SNAPSHOTCLASS    SNAPSHOTCONTENT
7 CREATIONTIME    AGE
8 rbd-pvc-snapshot    true            mysql-pv-claim    4Gi
9 csi-rbdplugin-snapclass    snapcontent-0f82a7cd-e84c-409e-80e0-6c54cf047379
10 131m            131m
11 [root@k8s-master01 examples]# kubectl delete volumesnapshot rbd-pvc-snapshot
12 volumesnapshot.snapshot.storage.k8s.io "rbd-pvc-snapshot" deleted
13
14 # 删除快照控制器
15 kubectl delete -f /root/k8s-ha-install/snapshotter/ -n kube-system

```

4.清理存储池

- 删除块存储池和文件存储池。

```

1 # 查看所有存储池
2 kubectl -n rook-ceph exec deploy/rook-ceph-tools -- ceph osd pool ls
3
4
5 kubectl delete -n rook-ceph cephblockpool replicapool
6 kubectl delete -n rook-ceph cephfilesystem myfs

```

5.清理 StorageClass

- 删除 Rook 创建的 StorageClass。

```

1 kubectl delete sc rook-ceph-block rook-cephfs
2
3 [root@k8s-master01 examples]# kubectl get sc
4 NAME                                PROVISIONER                                RECLAIMPOLICY    VOLUMEBINDINGMODE
5 ALLOWVOLUMEEXPANSION    AGE
6 rook-ceph-block    rook-ceph.rbd.csi.ceph.com    Delete            Immediate
7 true            38h
8 rook-cephfs        rook-ceph.cephfs.csi.ceph.com    Delete            Immediate
9 true            14h
10 [root@k8s-master01 examples]# kubectl delete sc rook-ceph-block rook-cephfs
11 storageclass.storage.k8s.io "rook-ceph-block" deleted
12 storageclass.storage.k8s.io "rook-cephfs" deleted

```

6.清理 Ceph 集群

- 删除 CephCluster 资源。

```

1 kubectl delete -f cluster.yaml

```

7.删除 Rook 资源

- 删除 Rook 的 Operator、Common 和 CRD 资源。

```
1 kubectl delete -f operator.yaml
2 kubectl delete -f common.yaml
3 kubectl delete -f crds.yaml
```

8.处理卡住资源（如有）

- 若资源删除卡住，参考 Rook 官方文档进行故障排除。如果由于某些原因操作员无法移除终结器（例如，操作员不再运行），您可以使用以下命令手动删除终结器：

```
1 for CRD in $(kubectl get crd -n rook-ceph | awk '/ceph.rook.io/ {print $1}'); do
2     kubectl get -n rook-ceph "$CRD" -o name | \
3     xargs -I {} kubectl patch {} --type merge -p '{"metadata":{"finalizers":
4     [null]}}'
```

几秒钟内，你应该能看到集群 CRD 已被删除，将不再阻止其他清理操作，例如删除 `rook-ceph` 命名空间。

如果命名空间仍然处于终止状态，您可以检查哪些资源正在阻止删除，并移除最终izers 并删除这些资源，

```
1 kubectl api-resources --verbs=list --namespaced -o name \
2 | xargs -n 1 kubectl get --show-kind --ignore-not-found -n rook-ceph
```

如果删除失败，终端一直卡在删除中那么应该是配置有 Finalizers 阻塞删除，Finalizer 是 Kubernetes 中一种机制，用于在资源删除前执行清理逻辑。如果 finalizer 未被释放，资源会处于删除挂起状态。可以通过 `kubectl get configmap rook-ceph-mon-endpoints -n rook-ceph -o yaml | grep finalizers` 确认了该 ConfigMap 存在 `finalizers`

```
[root@k8s-master01 examples]# kubectl api-resources --verbs=list --namespaced -o name \
> | xargs -n 1 kubectl get --show-kind --ignore-not-found -n rook-ceph
NAME                                DATADIRHOSTPATH  MONCOUNT  AGE  PHASE  MESSAGE  HEALTH  EXTERNAL  FSID
cephcluster.ceph.rook.io/rook-ceph /var/lib/rook    3          2d11h Deleting Deleting the CephCluster HEALTH_WARN dd96e25d-c610-4fc2-a0aa-41bf4cd9d750
cephfilesystemsubvolumegroup.ceph.rook.io/myfs-csi Ready myfs 23h
[root@k8s-master01 examples]# kubectl edit cepchcluster.ceph.rook.io/rook-ceph -n rook-ceph
cephcluster.ceph.rook.io/rook-ceph edited
[root@k8s-master01 examples]# kubectl api-resources --verbs=list --namespaced -o name | xargs -n 1 kubectl get --show-kind --ignore-not-found -n rook-ceph
error: the server doesn't have a resource type "cephclusters"
NAME                                PHASE  FILESYSTEM  QUOTA  AGE
cephfilesystemsubvolumegroup.ceph.rook.io/myfs-csi Ready myfs 23h
[root@k8s-master01 examples]# kubectl delete cepchfilesystemsubvolumegroup.ceph.rook.io/myfs-csi -n rook-ceph
cephfilesystemsubvolumegroup.ceph.rook.io "myfs-csi" deleted
```

如果删除过程中终端一直卡在这里没有删除成功，可能是 Finalizers 阻塞删除 那么可以使用 edit 删除字段

```
1 [root@k8s-master01 examples]# kubectl get cepchfilesystemsubvolumegroup.ceph.rook.io
myfs-csi -n rook-ceph -o yaml | grep finalizers
2     finalizers:
3
4
5 kubectl edit cepchfilesystemsubvolumegroup.ceph.rook.io myfs-csi -n rook-ceph
```

```

apiVersion: ceph.rook.io/v1
kind: CephFilesystemSubVolumeGroup
metadata:
  creationTimestamp: "2025-03-20T01:36:40Z"
  deletionGracePeriodSeconds: 0
  deletionTimestamp: "2025-03-20T16:34:49Z"
  finalizers:
    - cephfilesystemsubvolumeceph.rook.io
  generation: 3
  name: myfs-csi
  namespace: rook-ceph
  resourceVersion: "517988"
  uid: 9072aa85-0a8a-4a1b-aa2e-f34b485e32cc
spec:
  dataPoolName: ""
  filesystemName: myfs
  name: csi
  pinning:
    distributed: 1
status:
  info:
    clusterID: e1026845ad66577abae1d16671b464c8
    pinning: distributed=1
    observedGeneration: 2
    phase: Ready

```

删除这两行内容，即可删除

9.清理数据目录和磁盘

```
1 | rm -rf /var/lib/rook/*
```

10.清理OSD 所使用磁盘

```

1 | #!/usr/bin/env bash
2 | DISK="/dev/sdb"
3 |
4 | # Zap the disk to a fresh, usable state (zap-all is important, b/c MBR has to be clean)
5 |
6 | # You will have to run this step for all disks.
7 | sgdisk --zap-all $DISK
8 |
9 | # Clean hdds with dd
10 | dd if=/dev/zero of="$DISK" bs=1M count=100 oflag=direct,dsync
11 |
12 | # Clean disks such as ssd with blkdiscard instead of dd
13 | blkdiscard $DISK
14 |
15 | # These steps only have to be run once on each node
16 | # If rook sets up osds using ceph-volume, teardown leaves some devices mapped that lock
   the disks.
17 | ls /dev/mapper/ceph-* | xargs -I% -- dmsetup remove %
18 |
19 | # ceph-volume setup can leave ceph-<UUID> directories in /dev and /dev/mapper
   (unnecessary clutter)
20 | rm -rf /dev/ceph-*
21 | rm -rf /dev/mapper/ceph--*
22 |
23 | # Inform the OS of partition table changes
24 | partprobe $DISK

```

参考链接: <https://rook.io/docs/rook/v1.6/ceph-teardown.html#delete-the-dataon-hosts>